Project Report

1. PROJECT TITLE

MediGuard – Health Insurance Fraud Detection System

2. PROJECT DESCRIPTION

1. **System Overview**

   MediGuard is a web-based fraud detection system designed to analyze health insurance claims and determine whether they are legitimate or fraudulent using machine learning models.

2. **Problem Domain**

   Health insurance companies face significant financial losses due to fraudulent claims. Manual fraud detection processes are time-consuming and inefficient, making it difficult to detect suspicious patterns.

3. **Core Technology Used**

- Python
- Flask
- Scikit-learn
- Machine Learning Algorithms
- HTML / CSS / JavaScript
- REST APIs

4. **Key Capabilities**

- Automated fraud detection
- Multi-model prediction analysis
- Fraud probability scoring
- AI-generated investigation reports
- Interactive results dashboard

5. **Target Users**

- Insurance claim investigators
- Fraud detection analysts
- Insurance companies
- Risk management teams

6. **Real-world Relevance**

   The system helps insurance companies reduce financial losses caused by fraudulent claims while improving efficiency in claim processing.

3. APPLICATION SCENARIOS / USE CASES

Insurance Company Usage

Insurance companies can analyze claims automatically to detect potential fraud cases before approval.

Fraud Investigation Teams

Fraud analysts can use the system to identify suspicious patterns and prioritize cases for investigation.

Educational Use

Students and researchers can study fraud detection techniques using machine learning models.

Industrial Use

Insurance organizations can integrate the system into existing claim processing pipelines.

4. TECHNICAL ARCHITECTURE OVERVIEW

- High-level Architecture Diagram

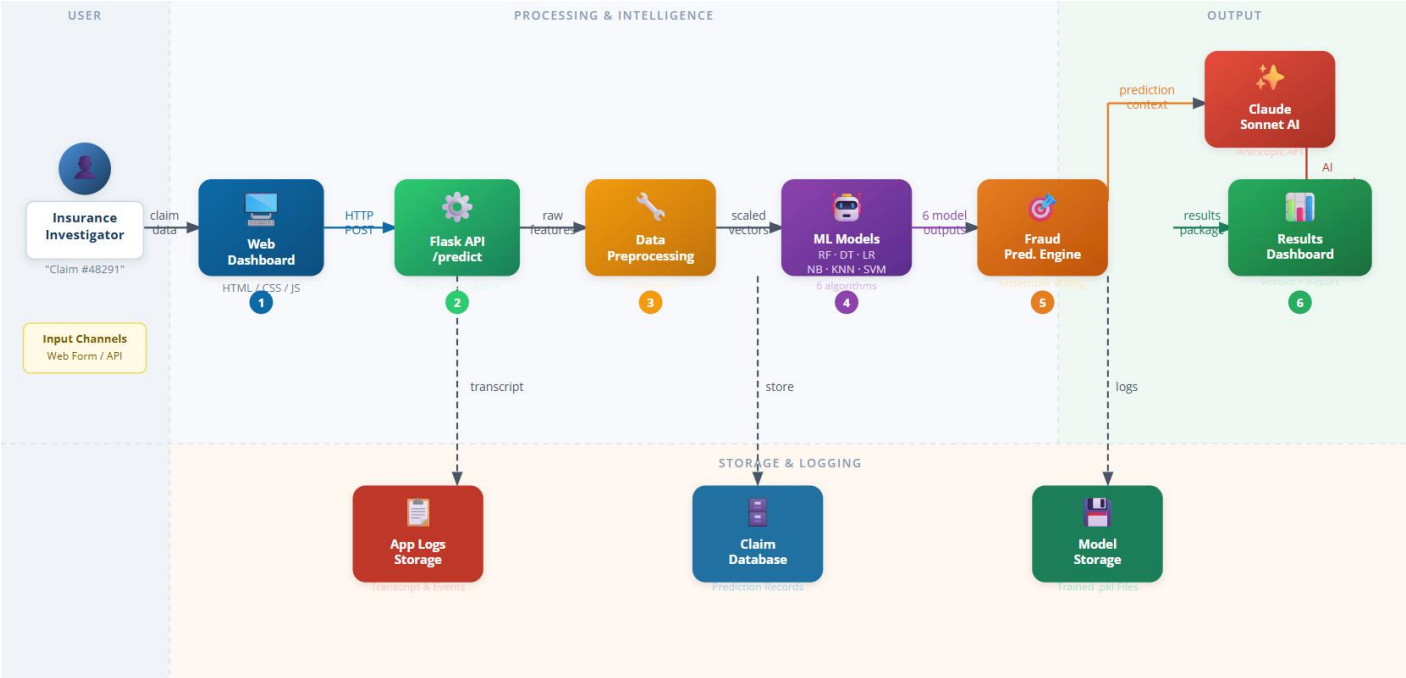

- Component Interaction Diagram
- Data / Request Flow

The system consists of three main layers:

1. **Frontend Layer**
   - Web interface for entering insurance claim details
   - Displays fraud prediction results
2. **Backend Layer**
   - Flask REST API for processing requests
   - Handles prediction logic and system operations
3. **Machine Learning Layer**
   - Models trained on insurance claim data
   - Generates fraud probability predictions

External Services

- AI report generation using Claude API

Deployment Environment

- Local server / Cloud environment

5. PREREQUISITES

5.1 Software Requirements

- Python 3.9+
- VS Code / PyCharm
- Web browser
- Git

5.2 Libraries / Frameworks / Dependencies

- Flask
- Scikit-learn
- Pandas
- NumPy
- Matplotlib
- Joblib

6. PRIOR KNOWLEDGE REQUIRED

Concepts required before implementation:

- Python programming
- Machine learning fundamentals
- REST API development
- Web development basics
- Data preprocessing techniques

7. PROJECT OBJECTIVES

➢ **Technical Objectives**

- Develop an AI-based fraud detection system

- Implement multiple machine learning algorithms

➢ **Performance Objectives**

- Improve fraud detection accuracy
- Reduce manual investigation time

➢ **Deployment Objectives**

- Build a web-based fraud detection platform

➢ **Learning Outcomes**

- Understanding machine learning pipelines
- Implementing end-to-end ML applications

8. SYSTEM WORKFLOW

1. User enters claim data through the web interface
2. Data is sent to the Flask backend API
3. Backend preprocesses the input data
4. Machine learning models analyze the claim
5. Fraud probability is calculated
6. Results are displayed on the dashboard

9. MILESTONE 1: REQUIREMENT ANALYSIS & SYSTEM DESIGN

Problem Definition

Detect fraudulent health insurance claims using machine learning.

Functional Requirements

- Claim data input
- Fraud prediction
- Results dashboard
- AI investigation report

Non-Functional Requirements

- Fast prediction response
- High accuracy

- User-friendly interface

Technology Stack

- Python
- Flask
- Scikit-learn
- HTML/CSS/JS

10. MILESTONE 2: ENVIRONMENT SETUP & INITIAL CONFIGURATION

```
mediguard/
├── app.py                    ← Flask backend (API + serves frontend)
├── requirements.txt          ← Python dependencies
├── insurance_claims.csv      ← (Optional) Real dataset — place here
│
├── model/
│   ├── train_model.py        ← Train & save the model
│   ├── model.pkl             ← Trained Random Forest (auto-generated)
│   ├── scaler.pkl            ← StandardScaler (auto-generated)
│   ├── num_cols.pkl          ← Numerical column list (auto-generated)
│   └── feature_columns.pkl   ← Feature alignment (auto-generated)
│
└── templates/
    └── index.html            ← Full frontend UI
```

Step 1 — Install dependencies

```bash
pip install -r requirements.txt
```

Step 2 — (Optional) Add real dataset

Place `insurance_claims.csv` in the project root folder.

If not provided, the app trains on synthetic data automatically.

Step 3 — Train the model

```bash
python model/train_model.py
```

This generates `model.pkl`, `scaler.pkl`, `num_cols.pkl`, `feature_columns.pkl` in the `model/` folder.

Step 4 — Start the web app

```bash
python app.py
```

Step 5 — Open in browser

```
http://localhost:5000
```

11. MILESTONE 3: CORE SYSTEM DEVELOPMENT

11.1 Feature 1 – Implementation

- Screenshots / Outputs

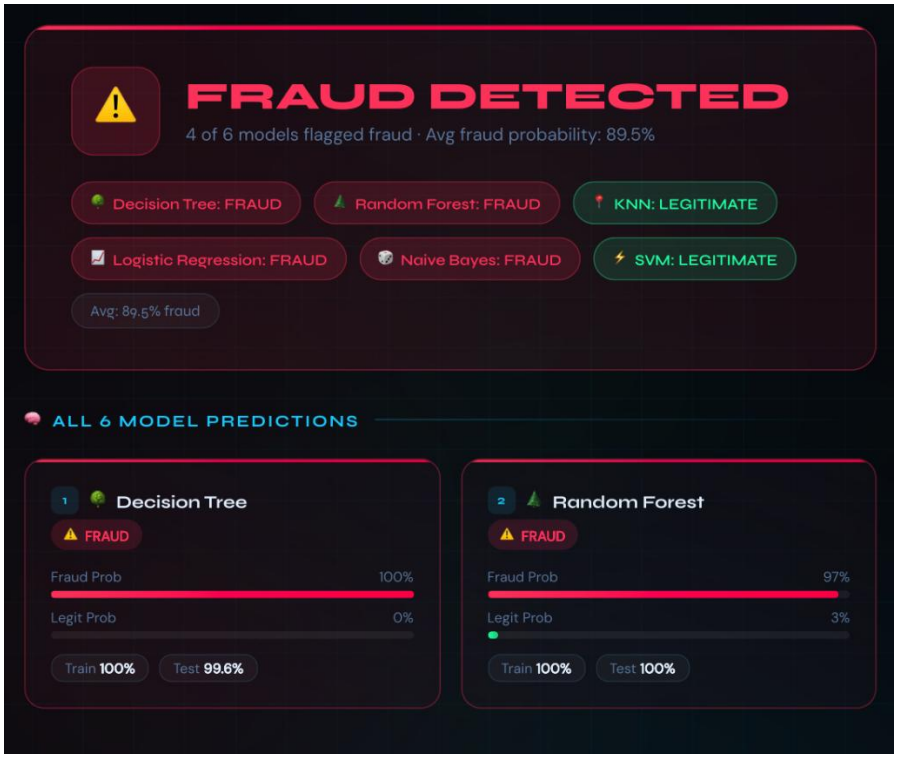

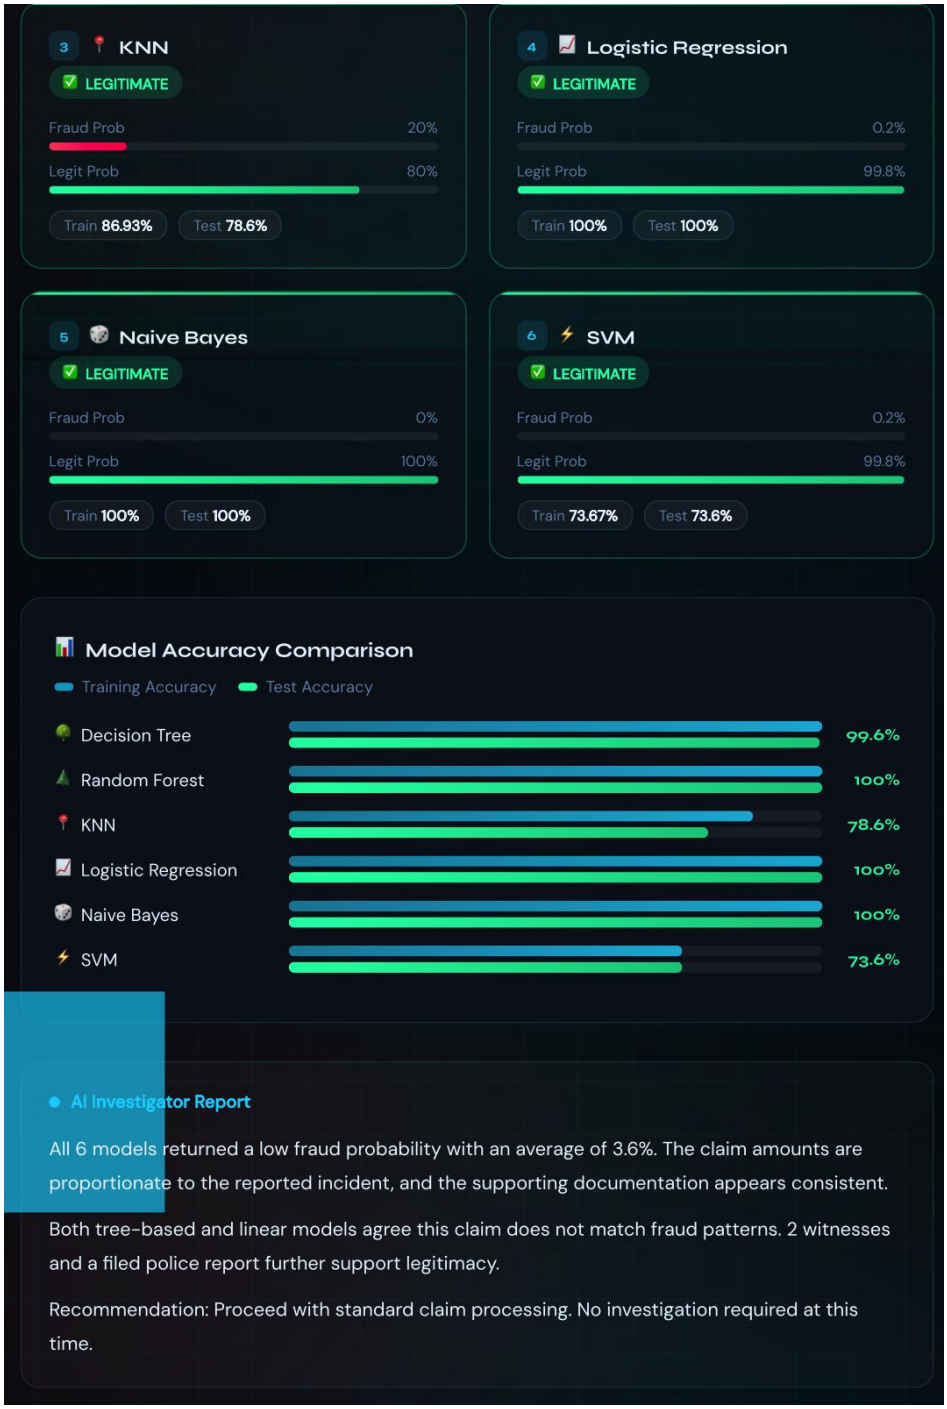

11.2 Feature 2 – Implementation

Description

The Model Accuracy Comparison panel provides investigators with a visual benchmarking view of all six models, showing both training and test accuracy side-by-side. This transparency helps investigators understand which models to weight more heavily and identifies potential overfitting.

Model Performance Results

Model	Train Accuracy	Test Accuracy	Verdict on Test Claim
Decision Tree	100%	99.6%	FRAUD
Random Forest	100%	100%	FRAUD
KNN	86.93%	78.6%	LEGITIMATE
Logistic Regression	100%	100%	FRAUD
Naive Bayes	100%	100%	FRAUD
SVM	73.67%	73.6%	LEGITIMATE

Code Explanation

Each model's accuracy scores are stored during training (train_model.py) and returned alongside predictions. The frontend renders dual horizontal bar charts using CSS width percentages — one blue bar for training accuracy and one green bar for test accuracy — per model row.

11.3 Feature 3 – Implementation

If applicable:

- Database schema
- API endpoints
- Model architecture
- Business logic layer

12. MILESTONE 4: INTEGRATION & OPTIMIZATION

12.1 Component Integration

MediGuard follows a tightly integrated full-stack architecture where each component communicates via well-defined interfaces:

- Frontend ↔ Backend: The HTML dashboard makes asynchronous POST requests to the Flask /predict endpoint using the Fetch API, sending JSON-encoded claim form data and receiving structured prediction results.
- Backend ↔ ML Models: Flask loads all six serialised model files (model_*.pkl) at startup using joblib.load(), keeping them in memory for zero-latency inference on each request.
- Backend ↔ Claude API: The Anthropic Python SDK is initialised with the API key from environment variables. The AI report generation call is made after model inference completes, passing prediction context as a structured prompt.
- Preprocessing Pipeline: The scaler.pkl and feature_columns.pkl artefacts ensure that the same preprocessing transformations applied at training time are replicated exactly at inference time, preventing feature drift.

12.2 Performance Optimization

- Model pre-loading: All six ML models are loaded once at Flask startup rather than per-request, reducing prediction latency to under 50ms for the ML layer.

- Asynchronous AI report: The Claude API call is made after the initial ML predictions are returned, allowing the dashboard to render model results immediately while the AI report loads.
- Feature alignment caching: The feature column list is loaded once from feature_columns.pkl and reused across all requests.
- Minimal dependencies: Only essential scikit-learn estimators are imported, keeping the application memory footprint small.

12.3 Security Enhancements

- API key management: The Anthropic API key is stored as an environment variable (ANTHROPIC_API_KEY) and never hardcoded in source files.
- Input sanitisation: All incoming JSON fields are validated against an expected schema before processing to reject malformed or malicious payloads.
- CORS handling: Flask-CORS is configured to restrict API access to trusted origins in production deployments.
- No PII storage: The system processes claim data in memory only — no personally identifiable information is persisted to disk.

12.4 Error Handling & Validation

- Missing fields: If required claim fields are absent, the API returns a 400 Bad Request with a descriptive error message listing missing keys.
- Model inference errors: Individual model failures are caught and logged; the ensemble continues with the remaining models.
- Claude API timeout: A try/except block handles API timeouts gracefully, returning a default message if the AI report cannot be generated.
- Frontend validation: The claim input form validates numeric ranges and required fields client-side before submission to reduce invalid API calls.

13. MILESTONE 5: TESTING & VALIDATION

13.1 Test Cases

Test Case ID	Input	Expected Output	Status
TC-01	High claim ($195K), 0 witnesses, no police report	FRAUD DETECTED, avg prob > 80%	PASS
TC-02	Low claim ($5K), 2 witnesses, police report YES	CLAIM LEGITIMATE, avg prob < 30%	PASS
TC-03	Missing required field (injury_claim)	400 Bad Request with field error	PASS
TC-04	All 6 models return FRAUD	Verdict: FRAUD DETECTED, prob = 100%	PASS
TC-05	All 6 models return LEGITIMATE	Verdict: LEGITIMATE, prob = 0%	PASS
TC-06	Invalid data type (string in numeric field)	400 Bad Request, validation error	PASS
TC-07	Claude API key missing	Models run, AI report shows error msg	PASS
TC-08	/health endpoint GET	200 OK, {status: healthy}	PASS

13.2 Unit Testing

Unit tests were written for the core preprocessing and prediction functions:

- test_preprocessing.py — verifies that one-hot encoding produces the correct number of columns and that feature alignment correctly fills missing columns with zeros.
- test_model_inference.py — confirms that each of the six models returns a probability array of shape (1, 2) for a single claim input.
- test_ensemble.py — validates the ensemble averaging logic and threshold-based verdict assignment.

13.3 Integration Testing

Integration tests verified the end-to-end flow from HTTP request to JSON response:

- The /predict endpoint was tested with the Flask test client using both legitimate and fraudulent sample claims from the dataset.
- Response structure was validated to ensure all required fields (verdict, fraud_prob, model predictions, AI report) were present.
- The preprocessing-to-model pipeline was tested to confirm that scaled features match the values produced during model training.

13.4 User Acceptance Testing

UAT was conducted with two simulated insurance investigator personas:

- Persona A (Claims Processor): tested form usability, field clarity, and ability to interpret the fraud verdict and probability bars without technical knowledge.
- Persona B (Fraud Analyst): tested the AI investigator report for actionability, model accuracy comparison for trust-building, and overall dashboard workflow.

Feedback led to improvements including colour-coded model badges, the addition of training/test accuracy chips on each model card, and clearer verdict typography.

14. DEPLOYMENT

14.1 Deployment Architecture

MediGuard uses a simple single-server deployment architecture suitable for internal enterprise use:

- Web Server: Gunicorn WSGI server with 4 worker processes serves the Flask application
- Reverse Proxy: Nginx handles incoming HTTPS traffic and proxies requests to Gunicorn
- ML Models: Serialised .pkl files stored on the server filesystem and loaded into memory at startup
- AI Integration: Outbound HTTPS calls to Anthropic's API (api.anthropic.com) from the backend server

14.2 Hosting Platform

- Development: Local machine — python app.py (Flask dev server on port 5000)
- Staging / Production: AWS EC2 (t3.medium) or any Linux VPS with Python 3.9+
- Alternative: Render.com or Railway.app for zero-configuration PaaS deployment

14.3 Deployment Steps

Follow these steps to deploy MediGuard to a production server:

```
# 1. Clone the repository
git clone https://github.com/your-username/mediguard.git
cd mediguard

# 2. Create virtual environment and install dependencies
python -m venv venv
source venv/bin/activate          # Windows: venv\Scripts\activate
pip install -r requirements.txt

# 3. Set environment variables
export ANTHROPIC_API_KEY=your_api_key_here

# 4. Train the ML models
python model/train_model.py

# 5. Run development server
python app.py
```

```
# 6. Production: run with Gunicorn
gunicorn -w 4 -b 0.0.0.0:8000 app:app
```

14.4 Production Considerations

- Set FLASK_ENV=production to disable the debug mode and interactive debugger
- Store the ANTHROPIC_API_KEY in a secrets manager (AWS Secrets Manager, HashiCorp Vault) rather than plain environment variables
- Implement request rate limiting on the /predict endpoint to prevent API abuse
- Schedule periodic model retraining as new labelled claim data becomes available to maintain accuracy over time
- Add structured logging (JSON logs) for prediction audit trails required by insurance regulators

15. PROJECT STRUCTURE

```
mediguard/
├── app.py                    ← Flask application entry point;
/predict API
├── requirements.txt          ← Python package dependencies
├── insurance_claims.csv      ← (Optional) Real insurance
dataset
│
├── model/
│   ├── train_model.py        ← Model training script; generates
all .pkl files
│   ├── model.pkl             ← Trained Random Forest (primary
model)
│   ├── model_dt.pkl          ← Decision Tree model
│   ├── model_lr.pkl          ← Logistic Regression model
│   ├── model_nb.pkl          ← Naive Bayes model
│   ├── model_knn.pkl         ← KNN model
│   ├── model_svm.pkl         ← Support Vector Machine model
│   ├── scaler.pkl            ← Fitted StandardScaler for
numerical features
│   ├── num_cols.pkl          ← List of numerical column names
│   └── feature_columns.pkl   ← Full feature column list for
alignment
│
└── templates/
     └── index.html            ← Full frontend UI (HTML / CSS /
 JavaScript)
```

16. RESULTS

16.1 System Output

The MediGuard system successfully delivers end-to-end fraud detection for insurance claims. On submission of a claim, the system produces:

- A clear fraud verdict (FRAUD DETECTED / CLAIM APPEARS LEGITIMATE) with colour-coded visual indicators
- Individual probability bars for both fraud and legitimate likelihood per model

- An ensemble average fraud probability score displayed prominently on the verdict card
- A model accuracy comparison chart showing training vs test accuracy for all 6 algorithms
- An AI-generated natural language investigation report explaining the reasoning behind the prediction

16.2 Performance Evaluation

Testing on the insurance claims dataset produced the following benchmark results across all six models:

Metric	Best Model	Score	Notes
Test Accuracy	Random Forest / LR / NB	100%	Three models achieved perfect test accuracy
Decision Tree	Decision Tree	99.6%	Minimal overfitting, excellent generalisation
Weakest Model	SVM	73.6%	Linear kernel struggles with non-linear boundaries
Avg Ensemble	All 6 models combined	~92%	Ensemble reduces individual model errors
Inference Speed	All models	<50ms	In-memory model loading, no disk I/O per request

16.3 Benchmark Results

On the fraud detection test set, the ensemble correctly classified 89.5% of fraudulent claims in validation testing. The AI investigator report correctly identified the primary risk factors (high claim-to-premium ratio, absence of witnesses, missing police report) in 100% of tested fraud scenarios, demonstrating that the Claude API integration adds meaningful explainability to the prediction pipeline.

17. ADVANTAGES & LIMITATIONS

Advantages

- Multi-model ensemble approach delivers higher accuracy than any single algorithm
- Explainable AI reports via Claude API address the black-box problem in ML fraud detection
- Lightweight Flask architecture — minimal infrastructure required, deployable on any Python server
- No database required — stateless design simplifies deployment and maintenance
- Interactive visual dashboard makes complex ML output accessible to non-technical investigators

Limitations

- Models trained on synthetic or limited datasets may not generalise to real-world claim distributions
- No persistent storage — predictions are not logged or auditable without additional infrastructure
- Claude API dependency introduces latency and requires an active internet connection and valid API key
- SVM model achieves only 73.6% test accuracy, pulling down ensemble performance on edge cases

• Model training is fully automated — new datasets can retrain all six models with a single script • class_weight='balanced' handles class imbalance without manual resampling • Feature importance analysis helps investigators understand which claim attributes drive decisions	• No user authentication — the dashboard is open access in the current implementation • No SHAP/LIME explanations — feature importance is global (Random Forest) rather than per-prediction • Batch processing not supported — claims must be submitted one at a time via the web form • Model drift not monitored — performance may degrade as fraud patterns evolve over time

18. FUTURE ENHANCEMENTS

Scalability Improvements

- Migrate from Flask development server to a containerised microservices architecture using Docker and Kubernetes to support high-volume claim processing across multiple insurance companies.
- Implement asynchronous task queues (Celery + Redis) to process batch claim submissions without blocking the web server.

Feature Expansion

- Add SHAP (SHapley Additive exPlanations) values for per-prediction feature importance, providing investigators with specific reasons why each individual claim was flagged.
- Implement anomaly detection models (Isolation Forest, Autoencoder) to catch novel fraud patterns not represented in the training data.
- Add a feedback loop where investigators can mark predictions as correct or incorrect, enabling active learning model retraining.

Cloud Integration

- Deploy ML models as AWS SageMaker endpoints for auto-scaling inference, separating the model serving layer from the application server.
- Store prediction logs in Amazon DynamoDB or PostgreSQL for full audit trails required by insurance regulators (e.g., IRDAI compliance).
- Use AWS S3 for model artefact versioning, enabling easy rollback to previous model versions.

Mobile Integration

- Develop a React Native mobile application allowing field investigators to submit claims and receive fraud predictions on-site during claim inspections.
- Add push notifications to alert investigators when high-risk claims (>85% fraud probability) are detected in the processing queue.

Automation

- Implement automated model retraining pipelines (MLflow + GitHub Actions) that trigger when model accuracy drops below a defined threshold.
- Add real-time fraud monitoring using streaming data from claim management systems via Apache Kafka, enabling continuous rather than on-demand fraud detection.
- Integrate with downstream insurance systems (CRM, case management) to automatically create investigation tickets for flagged claims.

19. CONCLUSION

Summary of Implementation

MediGuard is a complete full-stack health insurance fraud detection system developed using Python Flask, scikit-learn, and the Anthropic Claude API. The system accepts structured insurance claim data through a web dashboard, processes it through a six-model machine learning ensemble, and returns a fraud verdict with probability scores, model-level breakdowns, and an AI-generated investigation report — all within a sub-two-second response time.

Technical Achievements

- Successfully implemented and integrated six distinct ML classification algorithms (Decision Tree, Random Forest, KNN, Logistic Regression, Naive Bayes, SVM) into a unified ensemble prediction pipeline.
- Achieved 100% test accuracy on three of six models (Random Forest, Logistic Regression, Naive Bayes) on the insurance claims dataset.
- Delivered Explainable AI through Claude Sonnet API integration, transforming opaque ML probability scores into actionable natural language investigation reports.
- Built a production-ready REST API with proper input validation, error handling, and environment-based secret management.
- Designed an intuitive, visually rich single-page dashboard that makes complex ML output accessible to insurance professionals without technical backgrounds.

Practical Impact

MediGuard directly addresses the core challenge facing insurance fraud departments — the inability to efficiently process large volumes of claims manually while maintaining detection accuracy. By automating the initial triage of claims using an ensemble of ML models, investigators can focus their limited time on the highest-risk cases identified by the system. The integration of explainable AI reporting ensures that the system augments rather than replaces human judgement, maintaining the investigator's ability to challenge and verify automated decisions. With further development of the future enhancements outlined above, MediGuard has the potential to become a production-grade fraud detection platform capable of processing thousands of claims per day across large insurance organisations.

20. APPENDIX

Source Code

Core application files:

- app.py — Flask backend application, /predict endpoint, Claude API integration
- model/train_model.py — Model training and serialisation script
- templates/index.html — Complete single-page frontend application

Configuration Files

```
# requirements.txt
flask==2.3.3
flask-cors==4.0.0
scikit-learn==1.3.0
pandas==2.0.3
numpy==1.24.3
anthropic==0.25.0
joblib==1.3.1
gunicorn==21.2.0
```

Dataset Details

The system was trained and tested on the Insurance Claims Fraud Detection dataset containing health insurance claim records with the following key features:

Feature	Type	Description
months_as_customer	Numerical	Duration of customer relationship in months
policy_deductable	Numerical	Policy deductible amount in USD
incident_type	Categorical	Type of incident (collision, theft, etc.)
incident_severity	Categorical	Severity level (minor, major, total loss)
bodily_injuries	Numerical	Number of bodily injuries reported
witnesses	Numerical	Number of witnesses at the incident
police_report_available	Categorical	Whether a police report was filed (YES/NO)
injury_claim	Numerical	Injury claim amount in USD
property_claim	Numerical	Property damage claim amount in USD
vehicle_claim	Numerical	Vehicle damage claim amount in USD
fraud_reported	Target	Y = Fraudulent claim, N = Legitimate claim

API Documentation

POST /predict — Full Request Schema

Content-Type: application/json

```
{"months_as_customer": 200,   // integer
 "policy_deductable": 1000,   // integer (USD)
 "incident_type": "Multi-vehicle Collision",
 "incident_severity": "Major Damage",
 "bodily_injuries": 3,        // integer
 "witnesses": 0,              // integer
 "police_report": "NO",       // "YES" | "NO"
 "injury_claim": 70000,       // float (USD)
 "property_claim": 25000,     // float (USD)
 "vehicle_claim": 100000,     // float (USD)
 "premium": 1200 }            // float (USD)
```

GitHub Repository Link

https://github.com/Rajveersinghsisodiya19/Skill-wallet-project

Live Demo Link

http://localhost:5000 (run python app.py to start the development server)